

# ALEX CHEN

San Francisco, CA | (555) 382-9102 | alex.chen@email.com | linkedin.com/in/alexchen-ai | github.com/alexchen-ml

---

## PROFESSIONAL SUMMARY

Machine Learning Practitioner and Software Engineer with hands-on experience building, fine-tuning, and evaluating transformer-based NLP pipelines, tabular predictive models, and scalable REST APIs.

## EDUCATION

**University of California, Berkeley** — Bachelor of Science in Computer Science | Graduated: May 2024

Relevant Coursework: Natural Language Processing, Deep Learning, Distributed Systems, Database Systems

## TECHNICAL SKILLS

**Languages:** Python, SQL, C++, TypeScript, Bash

**Frameworks & Libraries:** PyTorch, Hugging Face Transformers, Scikit-learn, Pandas, NumPy, FastAPI, Flask

**Data & Cloud:** PostgreSQL, Redis, Docker, Git, AWS (S3, EC2 - familiar)

## WORK EXPERIENCE

**Applied AI Research Lab** — Machine Learning Intern (Berkeley, CA | June 2023 – August 2023)

- Built a transformer-based text classification system using PyTorch and Hugging Face to categorize 250,000+ scientific documents, achieving a 93.4% F1-score.
- Developed an automated data preprocessing and tokenization pipeline in Python, reducing model training preparation time by 38%.
- Benchmarked multiple embedding models against benchmark datasets and optimized batch inference latency from 140ms to 42ms per request.
- Collaborated with senior researchers to deploy an internal evaluation dashboard using FastAPI and Docker.

## PROJECTS

**Semantic Resume & Document Search Engine** | Python, PyTorch, FAISS, FastAPI

- Engineered a dense vector retrieval system converting resume documents and job requirements into 384-dimensional embeddings.
- Implemented FAISS cosine similarity indexing across 50,000 candidate profiles with sub-15ms retrieval latency.
- Created a RESTful API with FastAPI and integrated Redis for caching repeated queries, improving throughput by 45%.

**Financial Sentiment Analyzer** | Python, Scikit-learn, Pandas, Streamlit

- Developed an end-to-end sentiment classification model using Random Forest on 80,000 financial news headlines, reaching 89% accuracy.
- Deployed interactive web dashboard enabling real-time sentiment scoring for streaming market feeds.